

# Introduction to STM

## Part II: Transactional LEGOs

---

Laurens Duijvesteijn

June 2018

`duijf.io · github.com/duijf · @duijf`

## Preface

---

## Recap (1/2)

STM = Optimistic concurrency control

## Recap (1/2)

STM = Optimistic concurrency control

A TVar is a transactional variable

## Recap (1/2)

STM = Optimistic concurrency control

A **TVar** is a **transactional variable**

Operations on **TVars** are called STM transactions

## Recap (1/2)

STM = Optimistic concurrency control

A **TVar** is a **transactional variable**

Operations on **TVars** are called STM transactions

Run transactions with **atomically**; combine with **orElse**

## Recap (2/2)

Transactions get rerun on conflicts. We know because of a log

## Recap (2/2)

Transactions get rerun on conflicts. We know because of a log

You can also call **retry** to rerun yourself



## Recap (2/2)

Transactions get rerun on conflicts. We know because of a log

You can also call **retry** to rerun yourself

The runtime reruns the transaction after its inputs change

## Recap (2/2)

Transactions get rerun on conflicts. We know because of a log

You can also call **retry** to rerun yourself

The runtime reruns the transaction after its inputs change

Use this for blocking operations

## Today (by request)

Using `TVar` as a piece of LEGO

Let's create some of the datatypes in `stm`

If we have time, we can look at `stm-containers`

**Note:** this is all other people's code

## Transactional LEGOs

---

**TVars** are nice

But sometimes, we want something more

## **TVars** are nice

But sometimes, we want something more  
Certain situations call for more datatypes

## TVars are nice

But sometimes, we want something more

Certain situations call for more datatypes

The `stm` package provides a bunch of them

# Stuff in `stm`

**TVar** Variables



# Stuff in `stm`

**TVar** Variables

**TArray** Arrays (not discussing these)

## Stuff in `stm`

**TVar** Variables

**TArray** Arrays (not discussing these)

**TMVar** Variable which is either empty or full

## Stuff in `stm`

**TVar** Variables

**TArray** Arrays (not discussing these)

**TMVar** Variable which is either empty or full

**TQueue** Queues (bounded: **TBQueue**)

**TChan** Channels (slower than queues, but support broadcast)

TMVar

---

# Introducing TMVar a

Empty or filled (name comes from MVar, more later)

Can be used as a synchronization primitive

```
newTMVar :: a -> STM (TMVar a)
```

```
newEmptyTMVar :: STM (TMVar a)
```

```
putTMVar :: TMVar a -> a -> STM () -- blocks
```

```
takeTMVar :: TMVar a -> a -> STM a -- blocks
```

Can we build TMVar ourselves?

Yes.

# Can we build **TMVar** ourselves?

Yes. In Haskell, even.

# Can we build **TMVar** ourselves?

Yes. In Haskell, even. It's basically a **TVar** (Maybe a)



## Can we build `TMVar` ourselves?

Yes. In Haskell, even. It's basically a `TVar (Maybe a)`

**Really.** Add a `newtype` and that's the `stm` implementation

```
newtype TMVar a = TMVar (TVar (Maybe a))
```

```
newTMVar :: a -> STM (TMVar a)
```

```
newTMVar contents = do
```

```
    inner <- newTVar (Just contents)
```

```
    pure (TMVar inner)
```

```
newtype TMVar a = TMVar (TVar (Maybe a))
```

```
newEmptyTMVar :: STM (TMVar a)
```

```
newEmptyTMVar = do
```

```
    inner <- newTVar Nothing
```

```
    pure (TMVar inner)
```

```
putTMVar :: TMVar a -> a -> STM ()  
putTMVar (TMVar inner) newContents = do  
  contents <- readTVar inner  
  case contents of  
    Nothing -> do  
      writeTVar inner (Just newContents)  
      pure ()  
    Just _ -> retry
```

```
takeTMVar :: TMVar a -> STM a
takeTMVar (TMVar inner) = do
  contents <- readTVar inner
  case contents of
    Nothing -> retry
    Just c -> do
      writeTVar inner Nothing
      pure c
```

**MVar**

**TMVar**

## **MVar**

Defined in `Control.Concurrent.MVar`

## **TMVar**

Defined in `Control.Concurrent.STM.TMVar`

## **MVar**

Concurrent Haskell (1996)

## **TMVar**

Composable Memory Transactions (2006)



## **MVar**

(First come, first serve) fairness

## **TMVar**

**No** fairness guarantees

## **MVar**

Lower level tool. Can be used against starvation

## **TMVar**

High level convenient STM interface

# TQueue

---

## Introducing TQueue a

Unbounded, first-in first-out, queue

$O(1)$  inserts and reads (amortized).

```
newTQueue :: STM (TQueue a)
```

```
writeTQueue :: TQueue a -> a -> STM ()
```

```
readTQueue :: TQueue a -> STM a -- blocks
```

# Can we build `TQueue a` ourselves?

Again, yes. In this case using two `TVar [a]`s

## Can we build `TQueue` a ourselves?

Again, yes. In this case using two `TVar [a]`s

```
data TQueue a
  = TQueue (TVar [a]) -- waiting to be read
            (TVar [a]) -- written
```

```
data TQueue a = TQueue (TVar [a]) (TVar [a])
```

```
newTQueue :: STM (TQueue a)
```

```
newTQueue = do
```

```
    read  <- newTVar []
```

```
    write <- newTVar []
```

```
    pure (TQueue read write)
```

```
data TQueue a = TQueue (TVar [a]) (TVar [a])
```

```
writeTQueue :: TQueue a -> a -> STM ()  
writeTQueue (TQueue _read write) a = do  
    listend <- readTVar write  
    writeTVar write (a:listend)
```



```
data TQueue a = TQueue (TVar [a]) (TVar [a])
```

```
readTQueue :: TQueue a -> STM a
```

```
readTQueue (TQueue read write) = do
```

```
  xs <- readTVar read
```

```
  case xs of
```

```
    (x:xs') -> do
```

```
      writeTVar read xs'
```

```
      pure x
```

```
  [] -> do
```

```
    --
```

```
    -- What should go here??
```

```
    --
```

```
data TQueue a = TQueue (TVar [a]) (TVar [a])
```

```
readTQueue :: TQueue a -> STM a
```

```
readTQueue (TQueue read write) = do
```

```
  xs <- readTVar read
```

```
  case xs of
```

```
    (x:xs') -> do
```

```
      writeTVar read xs'
```

```
      pure x
```

```
  [] -> do
```

```
    -- Get the contents of `write`, reverse, and
```

```
    -- store in `read`. (retry if `write` is empty).
```

```
    -- This is where the "amortized" came from.
```

## Why add bounds?

Bounded queue: configurable (fixed) capacity. Writes block (**retry** with STM) if the new size is larger than the capacity.

Useful for implementing backpressure in our systems.

# How do we add bounds?

## How do we add bounds?

We need to know the sizes of the **read** and **write** lists.

## How do we add bounds?

We need to know the sizes of the **read** and **write** lists.

For every insert, we can see if we exceed capacity.

```
data TBQueue a
  = TBQueue (TVar Int) -- read capacity
              (TVar [a]) -- waiting to be read
              (TVar Int) -- write capacity
              (TVar [a]) -- written
```

```
newTBQueue :: Int -> STM (TBQueue a)
newTBQueue size = do
  read  <- newTVar []
  write <- newTVar []
  rsize <- newTVar 0
  wsize <- newTVar size
  pure (TBQueue rsize read wsize write)
```

```

writeTBQueue :: TBQueue a -> a -> STM ()
writeTBQueue (TBQueue rsize _read wsize write) a = do
  w <- readTVar wsize
  if (w /= 0)
    then do writeTVar wsize $! w - 1
    else do
      --
      --
      -- What should go here?
      --
      --
  listend <- readTVar write
  writeTVar write (a:listend)

```



```

writeTBQueue :: TBQueue a -> a -> STM ()
writeTBQueue (TBQueue rsize _read wsize write) a = do
  w <- readTVar wsize
  if (w /= 0)
    then do writeTVar wsize $! w - 1
    else do
      -- Figure out if we have read capacity.
      -- Retry if not.
      --
      -- If we have some, swap the read and
      -- write capacities and subtract 1.
  listend <- readTVar write
  writeTVar write (a:listend)

```

```
writeTBQueue :: TBQueue a -> a -> STM ()
writeTBQueue (TBQueue rsize _read wsize write) a = do
  w <- readTVar wsize
  if (w /= 0)
    then do writeTVar wsize $! w - 1
    else do
      r <- readTVar rsize
      if (r /= 0)
        then do writeTVar rsize 0
                writeTVar wsize $! r - 1
        else retry
  listend <- readTVar write
  writeTVar write (a:listend)
```

stm-containers

---

## Sometimes you want an STM (hash)map

This can be useful whenever a Hashmap is useful

Example: user-ID to rate limiting information.

A first idea might be `TVar (HashMap (TVar a))`, but that's not ideal.

## Why not TVar (HashMap (TVar a))?

Updates to individual keys are fine.

Updates to the entire **HashMap** change the outer **TVar**

This imposes a bottleneck for some workloads (contention on key addition/removal)

## We want to decrease contention

This requires splitting the workload into some chunks.

It does basically mean “implement **Map** again”.

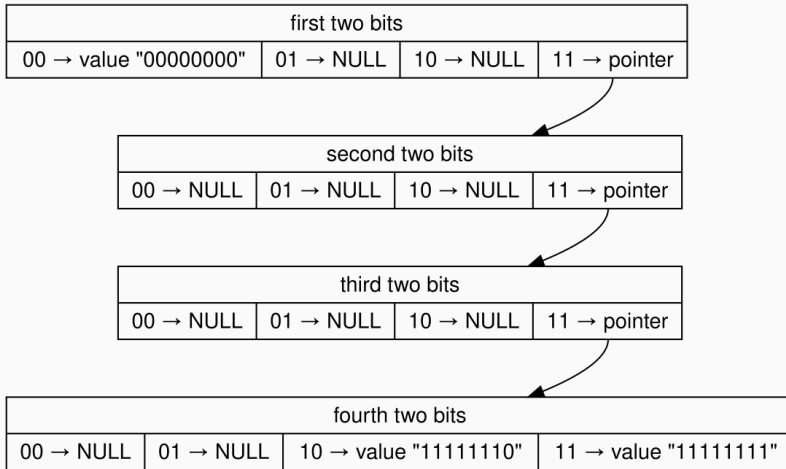
Method: a “Hash Array Mapped Trie”

## HAMT, briefly

Each key/value-pair is stored by:

1. Hashing the key (yielding a binary value)
2. Storing the value in the Trie at the location determined by the hash

More at: <https://idea.popcount.org/2012-07-25-introduction-to-hamt/>



Picture credit: <https://idea.popcount.org/2012-07-25-introduction-to-hamt/>



You can imagine this is better

# You can imagine this is better

Updates to a part of the Map are almost isolated

## You can imagine this is better

Updates to a part of the Map are almost isolated

We sometimes need to insert new levels, but this is waaay better than global contention.

## You can imagine this is better

Updates to a part of the Map are almost isolated

We sometimes need to insert new levels, but this is waaay better than global contention.

A real implementation is more involved and has compression and an bunch of other things.

More at: <https://nikita-volkov.github.io/stm-containers/>

## Summary

---

## Summary

`TMVar a = TVar (Maybe a)`

`TQueue a = TQueue (TVar [a]) (TVar [a])`

Want bounds? Add counts for read and write capacity!

**stm-containers** works with a Hash Array Mapped Trie to avoid contention on structural updates

# Tomorrow

First hour: excersises

Then: implementation notes and semantics